

SALSA20 : Implementation



Indian Statistical Institute, Kolkata

April 19th, 2024

Contents

- 1 Rise of Salsa20
 - Background
 - History
- 2 Overview
 - Agility of Salsa20
 - Speed
- 3 Salsa20 Specification
 - Initial State of *Salsa20*
 - The Quarterround function
 - Diagram
 - The Rowround function
 - The Columnround function
 - The Doublround function
- 4 Cryptanalysis on Salsa20
 - The Littleendian function
 - The Expansion function
 - The Hash function
 - The Encryption function
 - Brute Force Attacks
 - Hamming distance differential analysis
 - Differential analysis on QR function
 - Analysis on Salsa20's PRG
- 5 The Invincible Salsa20
- 6 Conclusion
- 7 References

Background

Background

- Use of network based applications are growing at a rapid speed.
- *Pseudo-random* numbers are at the core of any network security application.
- *Osvik, Shamir* and *Tromer* used cache-timing attacks to steal AES keys from a Linux disk-encryption device.
- Serious key collision & leakage in the hardware implementation of AES ciphers was found.
- *A. Shamir, I. Mantin* and *S. Fluhrer* revealed weaknesses in key scheduling algorithm of RC4.
- Cipher should be "**GENERIC**" compatible on both Hardware and Software platforms.

This way *Salsa20* came to the picture

Rise of Salsa20

History

- **eStream** : The *Ecrypt* Stream Cipher Project, called for submissions of stream ciphers in *November 2004*.
- **Salsa20** : Family of 256 – *bit* stream ciphers designed in 2005 and submitted to *eStream* by *Daniel J. Bernstein*.
- Salsa20 progressed to the *third round* of *eSTREAM* without any further changes.
- The final *eStream* portfolio, containing four software stream ciphers and four hardware stream ciphers, were announced in *April 2008*.
- It is not **patented**, and Bernstein has written several public domain implementations optimized for common architectures.

Overview

- Long chain of simple operations, rather than a shorter chain of complicated operations.
- This software-oriented stream cipher is built on a *Pseudorandom* function based on *ADD-ROTATE-XOR* (ARX) operations.
- It undergoes the following set of operations :
 - **32 – bit** Addition producing the sum $a + b \bmod 2^{32}$ of two 32 – bit words a , b .
 - **32 – bit** Exclusive-Or, producing the $a \oplus b$ of two 32 – bit words a , b .
 - Constant-distance **32**-bit rotation, producing the rotation $a \lll b$ of a 32 – bit word a by b bits to the **left** (where b is constant).

Agility of Salsa20

Arch	MHz	Machine	Salsa20 software	Cycles/byte					
				Salsa20/8		Salsa20/12		Salsa20/20	
				long	576	long	576	long	576
amd64	3000	Xeon 5160 (6f6)	amd64-xmm6	1.88	2.07	2.80	3.25	3.93	4.25
amd64	2137	Core 2 Duo (6f6)	amd64-xmm6	1.88	2.07	2.57	2.80	3.91	4.33
ppc32	533	PowerPC G4 7410	ppc-altivec	1.99	2.14	2.74	2.88	4.24	4.39
x86	2137	Core 2 Duo (6f6)	x86-xmm5	2.06	2.28	2.80	3.15	4.32	4.70
amd64	2000	Athlon 64 X2 (15,75,2)	amd64-3	3.47	3.65	4.86	5.04	7.64	7.84
ppc64	2000	PowerPC G5 970	ppc-altivec	3.28	3.48	4.83	4.87	7.82	8.04
amd64	2391	Opteron (f5a)	amd64-3	3.78	3.96	5.33	5.51	8.42	8.62
amd64	2192	Opteron (f58)	amd64-3	3.82	4.18	5.35	5.73	8.42	8.78
x86	2000	Athlon 64 X2 (15,75,2)	x86-1	4.50	4.78	6.27	6.55	9.80	10.07
x86	900	Athlon (622)	x86-athlon	4.61	4.84	6.44	6.65	10.04	10.24
ppc64	1452	POWER4	merged	6.83	7.00	8.35	8.51	11.29	11.47
hppa	1000	PA-RISC 8900	merged	5.82	5.97	7.68	7.85	11.39	11.56
amd64	3000	Pentium D (f64)	amd64-xmm6	5.38	5.87	7.19	7.84	10.69	11.73
x86	1300	Pentium M (695)	x86-xmm5	5.30	5.53	7.44	7.70	11.70	11.98
x86	3000	Xeon (f26)	x86-xmm5	5.30	5.86	7.41	8.21	11.64	12.55
x86	3200	Xeon (f25)	x86-xmm5	5.30	5.84	7.40	8.15	11.63	12.59
x86	2800	Xeon (f29)	x86-xmm5	5.33	5.95	7.44	8.20	11.67	12.65
x86	3000	Pentium 4 (f41)	x86-xmm5	5.76	6.92	8.12	9.33	11.84	13.40
x86	1400	Pentium III (6b1)	x86-mmx	6.37	6.79	8.88	9.29	13.88	14.29
sparc	1050	UltraSPARC IV	sparc	6.65	6.76	9.21	9.33	14.34	14.45
x86	3200	Pentium D (f47)	x86-athlon	7.13	7.66	9.90	10.31	15.29	15.94
ia64	1500	Itanium II	merged	8.49	8.87	12.42	12.62	18.07	18.27
ia64	1400	Itanium II	merged	8.28	8.65	12.56	12.76	18.21	18.40

Where we have :

$$\bullet \text{ Cycles / byte } = \frac{\text{cycles per Sec}}{\text{speed}}$$

$$\bullet \text{ speed } = \frac{\text{data size}}{\text{time}}$$

Figure: Speed on different platforms

Speed

- *Salsa20/20* runs at 3.93 cycles/byte for long streams. Whereas the fastest *AES* takes 9.2 cycles/byte for just 10 rounds of long stream.
- *Salsa20* runs at only 5.14 cycles/byte on a *Qualcomm Snapdragon S4 processor*, compared to 18.62 cycles/byte for *AES* – 128 in counter mode.
- 3 cycles/byte for *Cryptography* on *Core 2* *Salsa20/12* rounds takes 2.8 cycles/byte, one can afford at most 3 rounds of *AES* for any security at all.

Initial State of Salsa20

A **word** is an element of $\{0, 1, \dots, 2^{32} - 1\}$

The internal state is made of **sixteen** 32-bit words arranged in a 4×4 matrix. The initial state contains **eight** words of *key*, **two** words of *stream position*, **two** words of *nonce* (essentially additional stream position bits), and **four** *fixed words*:

"expa"	key	key	key
key	"nd 3"	nonce	nonce
block	block	"2-by"	key
key	key	key	"te k"

Table: *Salsa20's* initial state (**IS**) for 32 byte keys.

The Quarterround function

If x is a 4-word sequence then $quarterround(x)$ is also a 4-word sequence.

Definition

If $x = (x_0, x_1, x_2, x_3)$ then $quarterround(x) = (y_0, y_1, y_2, y_3)$ where :

$$y_1 = x_1 \oplus ((x_0 + x_3) \lll 7)$$

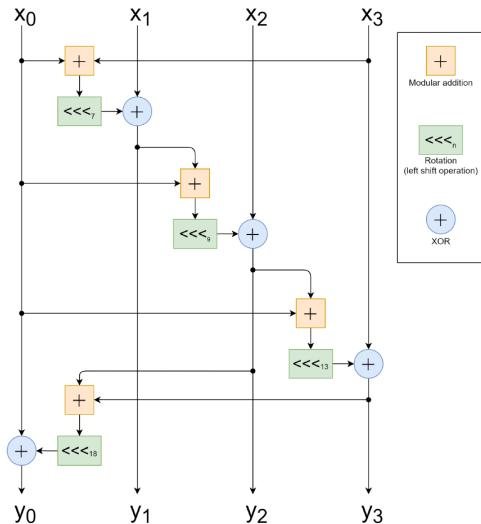
$$y_2 = x_2 \oplus ((x_1 + x_0) \lll 9)$$

$$y_3 = x_3 \oplus ((x_2 + x_1) \lll 13)$$

$$y_0 = x_0 \oplus ((x_3 + x_2) \lll 18)$$

N.B. : Each modification is **invertible**, so the entire function is **invertible**.

Diagram



`quarterround(0x00000001, 0x00000000, 0x00000000, 0x00000000) = (0x08008145, 0x00000080, 0x00010200, 0x20500000)`

The Rowround function

If y is a 16-word sequence then $\text{rowround}(y)$ is a 16-word sequence.

Definition

If $y = (y_0, y_1, y_2, y_3, \dots, y_{15})$ then
 $\text{rowround}(y) = (z_0, z_1, z_2, z_3, \dots, z_{15})$, where

$$(z_0, z_1, z_2, z_3) = \text{quarterround}(y_0, y_1, y_2, y_3)$$

$$(z_5, z_6, z_7, z_4) = \text{quarterround}(y_5, y_6, y_7, y_4)$$

$$(z_{10}, z_{11}, z_8, z_9) = \text{quarterround}(y_{10}, y_{11}, y_8, y_9)$$

$$(z_{15}, z_{12}, z_{13}, z_{14}) = \text{quarterround}(y_{15}, y_{12}, y_{13}, y_{14})$$

The Columnround function

If x is a 16-word sequence then $columnround(x)$ is a 16-word sequence.

Definition

If $x = (x_0, x_1, x_2, x_3, \dots, x_{15})$ then
 $columnround(x) = (y_0, y_1, y_2, y_3, \dots, y_{15})$ where,

$$(y_0, y_4, y_8, y_{12}) = quarterround(x_0, x_4, x_8, x_{12})$$

$$(y_5, y_9, y_{13}, y_1) = quarterround(x_5, x_9, x_{13}, x_1)$$

$$(y_{10}, y_{14}, y_2, y_6) = quarterround(x_{10}, x_{14}, x_2, x_6)$$

$$(y_{11}, y_{15}, y_3, y_7) = quarterround(x_{11}, x_{15}, x_3, x_7)$$

The Doubleround function

If x is a 16-word sequence then $\text{doubleround}(x)$ is a 16-word sequence.

Definition

A *doubleround* function is the composition of *columnround* followed by the *rowround* function, So we have

$$\text{doubleround}(x) = \text{rowround}(\text{columnround}(x))$$

The Littleendian function

If b is a 4-byte sequence then $littleendian(x)$ is a word.

Definition

If $b = (b_0, b_1, b_2, b_3)$ then we have,

$$littleendian(b) = b_0 + 2^8 \cdot b_1 + 2^{16} \cdot b_2 + 2^{24} \cdot b_3$$

Example

$$littleendian(255, 250, 126, 96) = 0x \underline{60} \underline{7e} \underline{fa} \underline{ff}$$

$$(255)_{28} = (\underline{1111} \underline{1111})_2 = (ff)_{16}$$

$$(250)_{28} = (\underline{1111} \underline{1100})_2 = (fa)_{16}$$

$$(126)_{28} = (\underline{0111} \underline{1110})_2 = (7e)_{16}$$

$$(96)_{28} = (\underline{0110} \underline{0000})_2 = (60)_{16}$$

The Salsa20 Expansion function

If k is a 32-byte and n is a 16-byte sequence then $Salsa20_k(n)$ is a 64-byte sequence.

Definition

Define $\sigma_0 = (101, 120, 112, 97)$, $\sigma_1 = (110, 100, 32, 51)$, $\sigma_2 = (50, 45, 98, 121)$, and $\sigma_3 = (116, 101, 32, 107)$. If k_0, k_1, n are 16-byte sequences then we have :

$$Salsa20_{k_0, k_1}(n) = Salsa20(\sigma_0, k_0, \sigma_1, n, \sigma_2, k_1, \sigma_3)$$

N.B. : Expansion refers to the expansion of (k, n) into $Salsa20_k(n)$. The constants $\sigma_0 \sigma_1 \sigma_2 \sigma_3$ is "expand 32 – byte k " in **ASCII**.

The *Salsa20* Hash function

If x is a 64-byte sequence then $Salsa20(x)$ is a 64-byte sequence.

Definition

$$Salsa20(x) = x + \text{doubleround}^{10}(x)$$

Where each 4-byte sequence is viewed as a word in *little – endian* form. Starting with $x = (x[0], x[1], \dots, x[63])$. Lets, define

$$x_0 = \text{littleendian}(x[0], x[1], x[2], x[3])$$

$$\vdots$$

$$x_{15} = \text{littleendian}(x[60], x[61], x[62], x[63])$$

Define $(z_0, z_1, \dots, z_{15}) = \text{doubleround}^{10}(x_0, x_1, \dots, x_{15})$

Then $Salsa20(x)$ is the concatenation of :

$$\text{littleendian}^{-1}(z_0 + x_0) \parallel \text{littleendian}^{-1}(z_1 + x_1) \parallel \dots \parallel \text{littleendian}^{-1}(z_{15} + x_{15})$$

The Salsa20 Encryption function

- Let k be a 32-byte sequence of key, v be a 8-byte sequence of nonce and m be an l -byte sequence of input, (for some $l \in \{0, 1, \dots, 2^{70}\}$).
- The *Salsa20 Encryption/Decryption of Input* is denoted by $\mathbf{Salsa20}_k(v) \oplus \mathbf{m}$, is an l -byte sequence.

Definition

$Salsa20_k(v)$ is a 2^{70} -byte sequence

$$Salsa20_k(v, 0) \parallel Salsa20_k(v, 1) \parallel \dots \parallel Salsa20_k(v, 2^{64} - 1)$$

$Salsa20_k(v) \oplus m$ implicitly truncates $Salsa20_k(v)$ to the same length as m . In other words : (where $c[i] = m[i] \oplus Salsa20_k(v, \lfloor \frac{i}{64} \rfloor)[i \bmod 64]$)

$$Salsa20_k(v) \oplus (m[0], m[1], \dots, m[l-1]) = (c[0], c[1], \dots, c[l-1])$$

Brute-force attacks

The Quarterround (QR) function takes a 128 bit binary number. The total possible combinations of inputs are 2^{128} .

A complete search would thus take about: (Assuming we can know the cryptographic nonce used)

$$\frac{2^{128} QR}{10000 QR/s} = 3.4 \cdot 10^{34} \text{ seconds} \approx 10^{27} \text{ years}$$

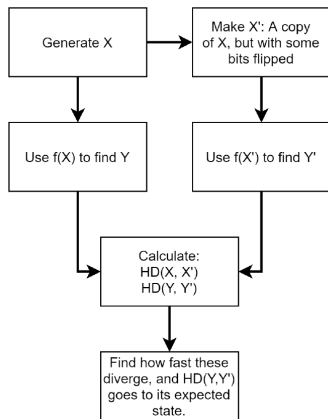
A complete search of the 256 bit key space would take:

$$\frac{2^{256} runs}{14 runs/s} = 2.4 \cdot 10^{76} \text{ seconds} \approx 10^{68} \text{ years}$$

Hamming distance differential analysis

For an *Encryption* function $f : X \rightarrow Y$

If $HD(X, X') = n$, then $HD(Y, Y') \stackrel{?}{=} m$



- If the algorithm has a good avalanche effect, we would expect the HD to be about the same as the HD between two random values: About half the bits.
- If P and Q are two random binary numbers of length n , we would expect: $HD(Q, P) \approx \frac{n}{2}$

Figure: Hamming Distance

Differential analysis on QR function

As $QR(X) \rightarrow y$ and we get X' by flipping some random bits of X . As n increases, m tend towards the expected equilibrium of half the length of Y . Analysis based on measuring $HD(QR(x'), y)$:

- The bit flipping is given by x -axis. i.e. n times.
- The HD between 2 values are given by y -axis.
- Legend of the Plot :
 - 0 $HD(X, Y)$
 - 1 Convergence point of the Hamming distance between two random values
 - 2 $HD(Y, Y')$
 - 3 $HD(X, X')$

Assuming $HD(X, X') < \frac{\text{key_size}}{8}$, $HD(X, X')$ should be easily distinguishable.




Figure: Flipping of random bits in X

Contd.




Figure: Averaging of effects on Y when bits are flipped in X .

Legend of the Plot :

- ① $HD(X', Y')$ as n bits in X' are flipped.
- ② $HD(X, Y)$ of random X s
- ③ Minimum expected difference between two random X s
- ④ Maximum expected difference between two random X s
- ⑤ Expected average difference between two random X s

The vertical line is where the bits flipped needed for the Hamming distance to be within the range of expected random distances

Analysis on Salsa20's PRG

- Line 1, shows how, $HD(input_{original}, input_{next})$ is roughly equal to 1 per flipped bits.
- Line 2, is the expected value for random inputs and outputs.
- Line 0, seems to be no correlation or pattern between the amounts of bits flipped in the $input(n)$ and the HD between the two values.




Figure: Averaging of effects on the PRG output when bits are flipped in key

The Invincible Salsa20

Salsa20 is highly resistant and secure against all the well known attacks :





- *Algebraic* attack
- *Weak — Key* attack
- *Equivalent — Key* attack
- *Related — Key* attack
- *Correlation power analysis*
- *Context aggregation network analysis*

Conclusion

After going through all this discussions, we conclude with the following points :

- *SALSA20* is faster and efficient as compared to *AES*.
- Been secure to both *KPA* and *CPA*.
- Efficient in both *software* and *hardware*.
- *Brute force* attack are not easily implementable.

References

-  D. J. Bernstein, “*The Salsa20 Family of Stream Ciphers*”, New Stream Cipher Des., pp. 84–97, 2008.
-  D. J. Bernstein, “*Salsa20 specification*”, eSTREAM Proj. algorithm Descr., pp. 2–10, 2005.
-  "Calculating cycles per byte." Stream cipher - Calculating cycles per byte - Cryptography Stack Exchange. N.p., 2 Oct. 2012. Web. 3 Mar. 2017.
<http://crypto.stackexchange.com/questions/3943/calculating-cycles-per-byte>
-  "How secure is Salsa20?" Algorithm design - How secure is Salsa20? - Cryptography Stack Exchange. N.p., 8 Oct. 2016. Web. 10 Mar. 2017.
<http://crypto.stackexchange.com/questions/40542/how-secure-is-salsa20/40543>